

Data Analytics Mini-Project (4-Day Sprint)

1. Project Overview

In this 4-day project, you will work **individually** with a real-world dataset from **Kaggle** to practice the full data analytics workflow:

- Finding and understanding a dataset
- Exploring, cleaning, and transforming data
- Building an analytical pipeline (descriptive analytics, simple models, or insights)
- Presenting your findings with **slides + live demo**

On **Friday**, each student will have:

- **10 minutes** for a slide presentation
- **10 minutes** for a live demo (notebooks / scripts / dashboard)

You are fully responsible for your own project, code, and presentation.

2. Learning Objectives

By the end of this project, you should be able to:

- Select a suitable dataset from Kaggle and understand its context.
- Design a **modular data pipeline** (exploration → cleaning → transformation → analytics).
- Apply good practices in **code organization**, **documentation**, and **version control** (Git + GitHub).

- Communicate technical work clearly through **visualizations, storytelling, and a demo**.
 - Follow a **style guide** and write clear, informative **docstrings**.
-

3. Timeline (4-Day Plan)

Day 1 – Project Setup & Dataset Selection

- Choose your Kaggle dataset (see Section 4 for constraints).
- Define your **project question(s)** and goals.
- Create a GitHub repository and basic project structure.
- Start initial EDA (Exploratory Data Analysis) notebook.

Day 2 – Cleaning & Transformation

- Deepen EDA (missing values, outliers, distributions, correlations, etc.).
- Design and implement a **cleaning pipeline** (in scripts/notebooks).
- Start defining transformation steps (feature engineering, encoding, scaling, etc.).

Day 3 – Analytics & Refinement

- Implement your Descriptive analytics and visualizations.
- Refine your pipeline and modularize code (functions, classes, scripts).
- Ensure code quality, docstrings, and style.

Day 4 – Final Polish & Presentation Prep

- Clean up notebooks (clear dead cells / unused code, add explanations and headings).
- Finalize repository structure and README.
- Prepare and rehearse your slide deck and demo.

- Make sure everything runs **from a clean environment** (reproducibility).

Friday – Presentations & Demos

- 10-minute slide presentation.
 - 10-minute live demo.
-

4. Dataset Selection (Kaggle)

You must choose a dataset from **Kaggle** that satisfies the following constraints:

- **Rows:**
 - At least **100,000 rows**
 - At most **1,000,000 rows**
- **Columns:**
 - At least **10 columns**
 - At most **30 columns**
- **Other requirements:**
 - Public, free-to-use Kaggle dataset.
 - Interesting to you (finance, health, sports, social media, environment, etc.).
 - Contains a mix of features (ideally some numeric and some categorical).

If the original dataset is larger than these limits, you may **download it and create a filtered/sample version** that meets the row/column constraints, but document clearly how you sampled it.

Required in your project:

- Include the **Kaggle link** in your README.

- Briefly explain in your README:
 - What is this dataset about?
 - Where does it come from?
 - What questions do you want to answer?
-

5. Project Requirements (Pipeline)

Your project must implement a clear pipeline with at least these stages:

1. Exploration

- Inspect data types, shapes, basic statistics.
- Explore missing values and distributions.
- Create meaningful visualizations (histograms, boxplots, scatter plots, bar charts, etc.).

2. Cleaning

- Handle missing values (drop, impute, domain-specific strategies).
- Handle outliers (remove, cap, or justify why you keep them).
- Fix data types and inconsistent values.

3. Transformation

- Encode categorical variables when needed.
- Scale or normalize features if relevant.
- Create new features (feature engineering) if it helps your analysis.

4. Analytics

- At minimum: strong descriptive analytics and clear visual storytelling.
- Optional but encouraged: a simple model (e.g., regression, classification, clustering) with basic evaluation metrics.

5. Conclusions

- Summarize key insights.
 - Discuss limitations and possible next steps.
-

6. Modularity & Project Structure

Your project must contain **modular functionalities**:

- Separate **notebooks** and **scripts** by purpose, for example:
 - `notebooks/01_eda.ipynb` – Initial exploration and visualizations
 - `notebooks/02_modeling.ipynb` – Modeling and evaluation
 - `src/cleaning.py` – Cleaning functions
 - `src/transformations.py` – Feature engineering & transformations
 - `src/utils.py` – Shared helper functions

Suggested directory structure:

```
project-root/
├── data/
│   ├── raw/      # Original data from Kaggle (do not modify)
│   └── processed/ # Cleaned / transformed datasets
├── notebooks/
│   ├── 01_eda.ipynb
│   └── 02_modeling.ipynb
└── src/
    ├── cleaning.py
    ├── transformations.py
    └── utils.py
└── reports/
    └── slides/    # Your final slide deck
└── README.md
└── requirements.txt # Python dependencies
```

7. Deliverables

By Friday, you must deliver:

1. GitHub Repository (per student)

- Clear, descriptive project name.
- Instructor has access.
- **Good README** (see Section 8).

2. Code + Notebooks

- Modular Python scripts in `src/`.
- Notebooks in `notebooks/` with:
 - Clean, readable structure (headings, explanations).
 - Minimal unused or commented-out code.
- All code must run without errors from a fresh clone (assuming correct environment).

3. Slides

- Final slide deck in `reports/slides/`.
- Follows the template in Section 11 (you can adapt, but keep the core structure).

4. Live Demo

- Show how your repository is organized.
- Show how to run at least one key notebook or script.

8. README Requirements

Your `README.md` should include at least:

- **Project Title**
- **Short Description** (2–3 sentences)
- **Dataset**
 - Link to Kaggle dataset
 - Short description and motivation
- **Dataset Size**
 - Number of rows and columns used in your project
 - Mention if you sampled or filtered from a larger original dataset
- **Project Goals / Questions**
 - What are you trying to understand or predict?
- **Project Structure**
 - Brief explanation of folders and main scripts/notebooks.
- **How to Run**
 - Environment requirements (Python version, main libraries).
 - How to install dependencies (`pip install -r requirements.txt`).
 - How to run the main analysis (which notebooks/scripts to execute).
- **Results & Key Insights**
 - 3–5 bullet points summarizing your main findings.

9. Code Quality & Style Guide

We will review **code quality and style**. You are expected to:

- Use **meaningful names** for variables, functions, and classes.
- Follow a consistent style (e.g., PEP 8 for Python):
 - snake_case for functions and variables (e.g., `clean_missing_values`).
 - PascalCase for classes (e.g., `DataCleaner`).
- Avoid very long functions: break logic into smaller, reusable functions.
- Remove dead code and unnecessary print statements.
- Keep notebooks clean and narrative: mix code with explanations and headings.

10. Documentation & Docstrings

Every **function** and **class** in your `src/` folder must include a docstring with:

- **Short description** (one sentence)
- **Long description** (optional, 2–4 sentences if useful)
- **Args** (list of parameters, types, and meaning)
- **Returns** (type and meaning of the output)
- **Raises** (exceptions that the function may raise, if any)

Example (Google-style docstring):

```
def clean_missing_values(df, strategy="mean"):
    """
    Clean missing values in a DataFrame.

    This function handles missing values according to the chosen
    strategy.

    Currently supported strategies are "mean", "median", and "drop".

    Args:
        df (pandas.DataFrame): Input DataFrame with missing values.
        strategy (str): Strategy to handle missing values. One of
            {"mean", "median", "drop"}.

    Returns:
        pandas.DataFrame: A new DataFrame with missing values handled
            according to the specified strategy.

    Raises:
        ValueError: If an unsupported strategy is provided.
    """
    ...

```

You may choose Google-style, NumPy-style, or similar, but be **consistent** across the project.

11. Presentation Guidelines (Slides + Demo)

Each student has **20 minutes total**:

- **10 minutes – Slide presentation**
- **10 minutes – Live demo**

Suggested Slide Template (10–12 slides)

1. Title Slide

- Project title
- Your name

2. Motivation & Problem Statement

- Why did you choose this dataset?
- What questions are you trying to answer?

3. Dataset Overview

- Kaggle dataset link
- Rows, columns (and confirm it meets the project constraints)
- Main features and context

4. Data Exploration Highlights

- Key descriptive statistics
- Interesting distributions or correlations
- 2–4 key plots with brief interpretation

5. Data Cleaning

- Main data issues (missing values, outliers, inconsistencies)
- How you handled them (drop, impute, transform, etc.)

6. Transformations & Feature Engineering

- Encodings, scaling, or feature creation
- Rationale for your choices

7. Analytics / Modeling

- What types of analyses or models did you apply?
- Key metrics or findings (accuracy, RMSE, clusters, patterns, etc.)

8. Results & Insights

- 3–5 key takeaways
- What did you learn from the data?

9. Limitations

- What are the limitations of your data or approach?
- What would you improve with more time?

10. Conclusion & Next Steps

- Summary of the project
- Ideas for future work or extensions

11. Live Demo Intro (optional)

- A quick slide explaining what you will show in the demo.

During the **demo**, show:

- How your repository is organized.
 - How to run at least one key notebook or script.
 - A highlight of one or two important visualizations or results.
-

12. Evaluation Criteria (High-Level)

You will be evaluated on:

1. Technical Work (Pipeline)

- Quality of exploration, cleaning, transformation, and analytics.
- Correct and meaningful use of methods.

2. Code Quality & Structure

- Modularity, readability, style, and organization.
- Proper use of functions, classes, and docstrings.

3. Reproducibility

- Clear instructions in README.
- Code runs without errors on a fresh clone (within reasonable setup).

4. Analysis & Insight

- Depth of understanding of the data.
- Clarity and relevance of insights / conclusions.

5. Communication

- Slide design, structure, and clarity.
 - Quality of explanations during the presentation and demo.
 - Ability to answer questions about your pipeline and choices.
-

13. Collaboration & Academic Integrity

This is an **individual** project.

- You may:
 - Discuss high-level ideas or challenges with classmates.
 - Use online resources (documentation, tutorials, etc.).
- You must not:
 - Copy another student's code or notebooks.
 - Present external work as if it were entirely your own.

Always:

- Cite any tutorials, blogs, or code snippets you used or adapted (mention them in README or notebook comments).

14. Checklist Before Friday

- Kaggle dataset chosen that satisfies size constraints (100k–1M rows, 10–30 columns).
- GitHub repo created and shared with instructor.
- Project structure organized (data/, notebooks/, src/, reports/).
- Cleaning and transformation scripts implemented with docstrings.
- At least one EDA notebook and one analytics/modeling notebook.
- README completed with how-to-run instructions and key insights.
- Slides finished following the template.
- Demo tested: can you run key steps from scratch?